

Michael Burke, Billy Theisen

7 April 2017

Project 5

Purpose

The purpose of the project is to understand how virtual memory management works and to create a few page replacement algorithms. As a group, we wrote a page fault handler that moves data from disk onto physical memory or sets the appropriate read and write bits when a page fault occurs. We ran the experiment of student machine 03 and the command line arguments used to test were:

Random:

```
./virtmem 100 3 rand focus
```

```
./virtmem 100 10 rand focus
```

...

```
./virtmem 100 100 rand focus
```

Custom:

```
./virtmem 100 3 custom focus
```

```
./virtmem 100 10 custom focus
```

...

```
./virtmem 100 100 custom focus
```

FIFO:

```
./virtmem 100 3 fifo focus
```

```
./virtmem 100 10 fifo focus
```

...

`./virtmem 100 100 fifo focus`

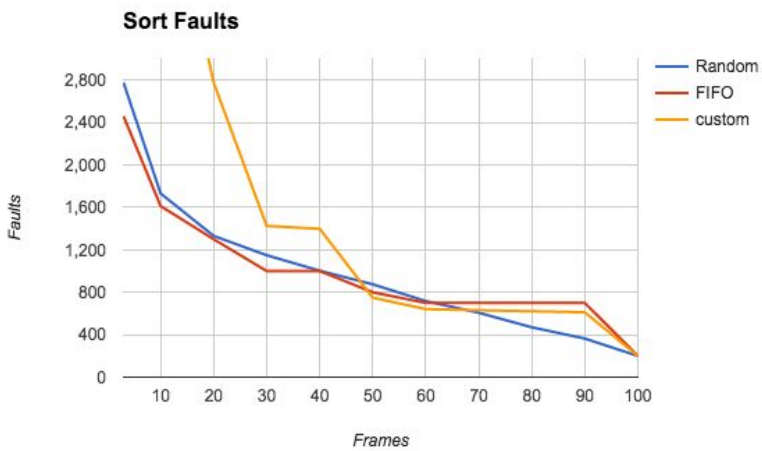
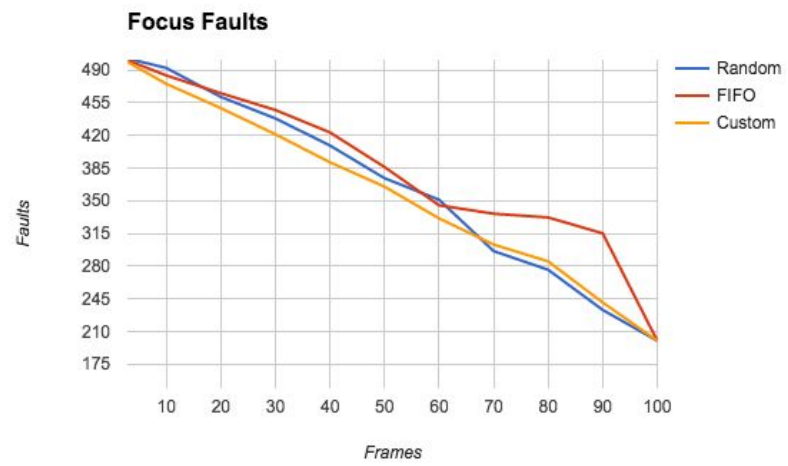
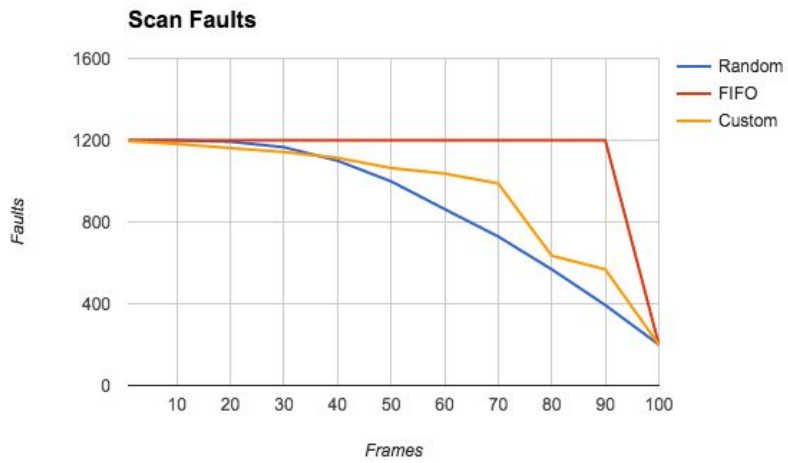
These arguments were changed to handle scan and sort as well.


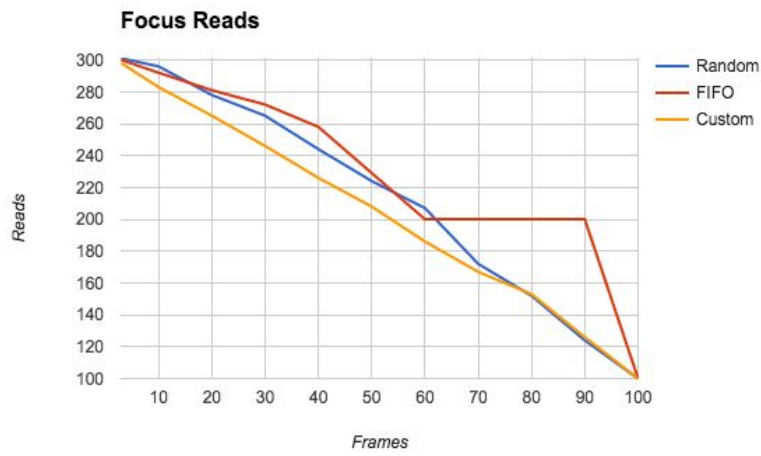
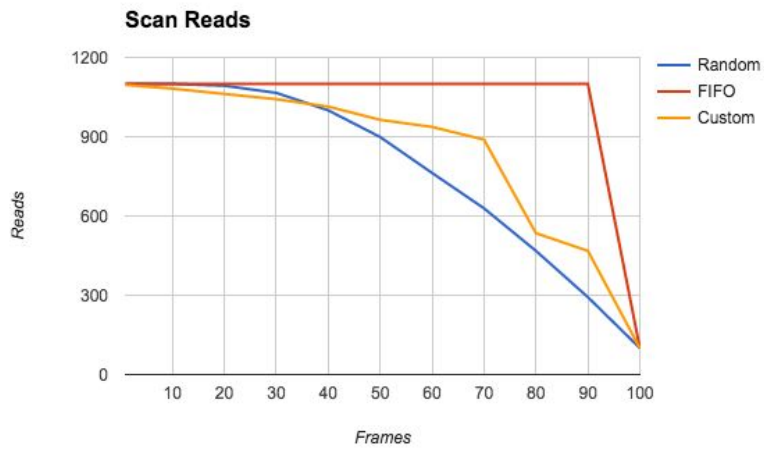
Custom Algorithm

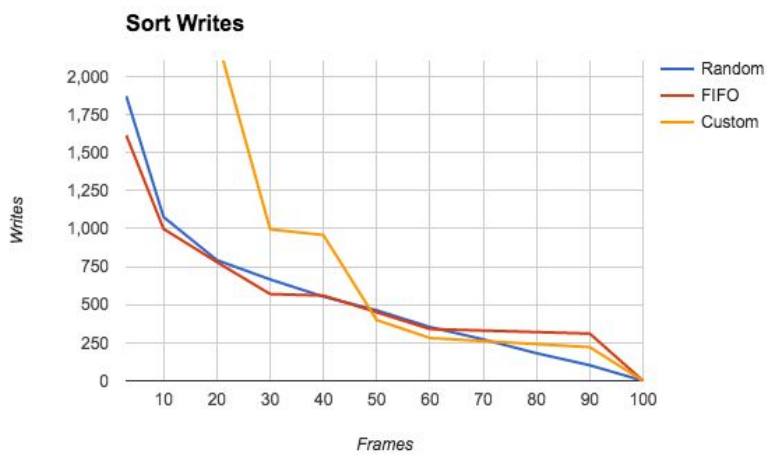
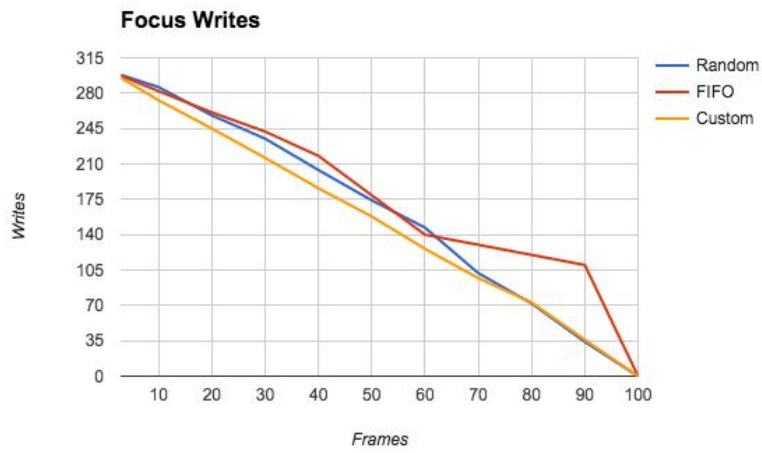

Our custom algorithm replaces a page that was written to the longest time from the page fault. Using a struct called `timeval`, once a page has its written bits set, we store a timestamp to an array called `accessTable->elements[]`. This timestamp struct has two values in it: seconds and its remaining microseconds. Pages that are not currently loaded into physical memory or do not have its written bits set have timestamps of 0.

When replacement occurs, we iterate over the `accessTable->elements[]` finding the minimum value which represents the oldest written page. We compare each element's timestamp second values. If the seconds equal each other, we then compare the microseconds. We give a priority boost to recently written pages. We add 10,000 microseconds to these pages' timestamp. Handling this addition is tricky. When we add 10,000 to the microseconds variable, sometimes the resultant variable is greater than or equal to one second. When this occurs, we need to instead add one to the seconds variable and then subtract 900,000 from microseconds variable which is a total time increase of 10,000 microseconds.

Results







Analysis

For the program focus, all three algorithms performed at about the same rate, with our custom algorithm slightly beating random and FIFO in faults, reads, and writes. Custom's technique of replacing the oldest written page must give it a slight advantage against random and FIFO.

For the scan program, random performs the best while FIFO performs the worst.

For the sort program, custom performs very very poorly with less than 10 or 20 frames. Once the frames equal 30 or more, custom performs just as well as random and FIFO. The branching of quicksort's recursion or the divide and conquer nature of the program must make custom very inefficient with such little frames as 10 and 20.

FIFO has an intrinsic flaw in its logic where it is inefficient in programs that in order access data.

This may be why FIFO's performance lags behind random and custom against all three programs.